

Convergent Cross Mapping: does A actually drive B?

Leonardo Capitani

2026-07-22

```
library(rEDM)
library(tidyverse)
```

This is the second of two notebooks. The [first one](#) covered reconstructing a *single* variable's dynamics with Simplex and S-map. Here I ask a different question: given two time series, can I tell whether one actually drives the other?

The question CCM answers

- You've got two time series — say, a nutrient and an algal population, or a predator and its prey.
- They're correlated. **Q: does that mean one causes the other?**
- A: Not necessarily — and the usual fixes (Granger causality, adding lags in a regression) assume the variables are *separable*: that you could, in principle, isolate X 's effect from everything else going on.
- In a tightly coupled ecological system, that assumption is often false. **Convergent Cross Mapping (CCM)** was built specifically for non-separable, dynamically coupled variables (Chang *et al.* 2017).

Does CCM need to control for other variables?

This is a real question people ask, and the honest answer has two parts.

Part 1 — no, not the way regression does.

- Multiple regression controls for confounders by explicitly adding them as covariates, because it treats each variable's contribution as separable and additive.

- CCM doesn't need that: if two variables are part of the same dynamical system, Takens' theorem guarantees that each variable's own lagged reconstruction already contains a "shadow" of everything it interacts with — including variables you never measured.
- So you're not omitting a confounder by leaving it out of the embedding; its influence is already folded into the variable you did measure.

Part 2 — but CCM has its own failure modes, and they need their own fixes.

- **Synchrony from strong one-way forcing.** If X drives Y *very* strongly, Y 's dynamics can collapse onto X 's so tightly that cross mapping looks bidirectional even though the true influence only runs one way.
- **Shared external drivers.** If X and Y don't interact at all, but both are pushed by the same external cycle (e.g. season), they can still cross-map each other successfully — a real confound, just not one you fix by "adding a covariate."
- Ye *et al.* (2015) give you the actual fix for both: look at cross-map skill **as a function of an added time lag**, not just at lag zero. I'll do exactly that below.

The core mechanic: cross mapping

- Reconstruct each variable's own manifold (same lagged-embedding idea as [the previous notebook](#)).
- If X causes Y , then information about X gets written into Y 's dynamics — so Y 's manifold should be able to "recover" X .
- **The direction rule (the confusing part):** to ask "*does X cause Y ?*", you cross-map using **Y 's manifold to predict X** — not the other way round. High, convergent skill there is your evidence for $X \rightarrow Y$.
- **Convergence** is the other half of the name: as you give the cross-mapping more library data (longer time series), skill should keep improving if the link is real. A spurious link (e.g. two unrelated series that happen to correlate) won't converge.

Step 1 — a sanity check where I already know the answer

Before trusting CCM on real data, I want to see it recover a causal link I built in myself. I'll simulate a nutrient pool driving an algal population strongly, with almost no feedback the other way:

```
set.seed(42)
n <- 400
b_algae_on_nutrients <- 0.02 # algae barely affects nutrients
b_nutrients_on_algae <- 0.32 # nutrients strongly drive algae -> TRUE causal link: Nutrient
Nutrients <- numeric(n); Algae <- numeric(n)
```

```

Nutrients[1] <- 0.4; Algae[1] <- 0.2

for (t in 2:n) {
  Nutrients[t] <- Nutrients[t-1] * (3.8 - 3.8 * Nutrients[t-1] - b_algae_on_nutrients * Algae[t-1])
  Algae[t] <- Algae[t-1] * (3.5 - 3.5 * Algae[t-1] - b_nutrients_on_algae * Nutrients[t-1])
  Nutrients[t] <- pmin(pmax(Nutrients[t], 1e-6), 0.9999)
  Algae[t] <- pmin(pmax(Algae[t], 1e-6), 0.9999)
}

sim_data <- tibble(time = 1:n, Nutrients = Nutrients, Algae = Algae)

```

Now I run CCM in both directions and watch whether skill *converges* as I hand it more data:


```

ccm_sim <- CCM(
  dataFrame = sim_data, E = 3, columns = "Nutrients", target = "Algae",
  libSizes = "10 300 20", sample = 100, seed = 1, showPlot = FALSE
)

ccm_sim |>
  pivot_longer(-LibSize, names_to = "direction", values_to = "rho") |>
  mutate(direction = recode(direction,
    "Algae:Nutrients" = "Algae's manifold predicts Nutrients -> evidence Nutrients causes Algae",
    "Nutrients:Algae" = "Nutrients's manifold predicts Algae -> evidence Algae causes Nutrients",
  )) |>
  ggplot(aes(x = LibSize, y = rho, color = direction)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2) +
  scale_color_manual(values = c("#1B7837", "grey60")) +
  labs(x = "Library size (L)", y = expression(rho), color = NULL,
    title = "CCM recovers the causal direction I built in") +
  theme_bw(base_size = 14) +
  theme(legend.position = "bottom")

```

CCM recovers the causal direction I built in



Nutrients → evidence Nutrients causes Algae —●— Nutrients's manifold predicts Al

Q: which curve should win, and does it? A: “Algae’s manifold predicts Nutrients” should climb and convergence — because Nutrients truly drives Algae — while the reverse should stay flat and low. That’s what the plot shows. Good: CCM recovers a direction I already know is right.

Step 2 — real predator-prey data

Now the real test: a classic laboratory microcosm, *Paramecium* (prey) and *Didinium* (its predator), the same dataset used by Ye *et al.* (2015).

```
data(paramecium_didinium)
paramecium_didinium |> head()
```


	time	paramecium	didinium
1	0.00	15.65	5.76
2	0.52	53.57	9.05
3	1.01	73.34	17.26
4	1.54	93.93	41.97
5	2.04	115.40	55.97
6	2.51	76.57	74.91

```

ccm_real <- CCM(
  dataFrame = paramecium_didinium, E = 3, columns = "paramecium", target = "didinium",
  libSizes = "10 65 5", sample = 200, seed = 1, showPlot = FALSE
)

ccm_real |>
  pivot_longer(-LibSize, names_to = "direction", values_to = "rho") |>
  mutate(direction = recode(direction,
    "didinium:paramecium" = "paramecium -> didinium (prey drives predator)",
    "paramecium:didinium" = "didinium -> paramecium (predator drives prey)"
  )) |>
  ggplot(aes(x = LibSize, y = rho, color = direction)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2) +
  scale_color_manual(values = c("#1B7837", "#762A83")) +
  labs(x = "Library size (L)", y = expression(rho), color = NULL,
    title = "Paramecium & Didinium: who's driving whom?") +
  theme_bw(base_size = 14) +
  theme(legend.position = "bottom")

```



ium -> paramecium (predator drives prey) — paramecium -> didinium (prey drives predator)

Q: what does it mean if *both* curves converge? A: Real, reciprocal causation — exactly what a predator-prey pair should look like: prey abundance drives predator growth (more food, more offspring) *and* predator abundance drives prey decline (more mouths, fewer survivors).

Unlike the synthetic example, this is genuinely bidirectional, and both curves climbing is the correct result, not a red flag.

Is it real, or could 65 points show this by chance?

Convergence is reassuring, but I can also ask directly: is the observed cross-map skill higher than I'd get from a version of the data with the same basic statistics but no real dynamic link? I scramble one series (destroying its dynamics, keeping its distribution) and rerun CCM many times to build a null distribution:

```
set.seed(1)
n_surr <- 100


surrogates <- SurrogateData(paramecium_didinium$didinium, method = "random_shuffle", num_surr = n_surr)

null_rho <- map_dbl(seq_len(n_surr), function(i) {
  d <- paramecium_didinium
  d$didinium <- surrogates[, i]
  out <- CCM(dataFrame = d, E = 3, columns = "paramecium", target = "didinium",
    libSizes = "65 65 1", sample = 100, seed = 1, showPlot = FALSE)
  out$`didinium:paramecium`[1]
})

observed_rho <- ccm_real$`didinium:paramecium`[ccm_real$LibSize == max(ccm_real$LibSize)]

tibble(rho = null_rho) |>
  ggplot(aes(x = rho)) +
  geom_histogram(bins = 20, fill = "grey80", color = "white") +
  geom_vline(xintercept = observed_rho, color = "#1B7837", linewidth = 1) +
  annotate("text", x = observed_rho, y = Inf, label = "observed", vjust = 2,
    hjust = -0.1, color = "#1B7837") +
  labs(x = expression(rho~"(shuffled surrogates)"), y = "Count",
    title = "Observed skill vs. a shuffled-data null distribution") +
  theme_bw(base_size = 14)
```

Observed skill vs. a shuffled-data null distribution



```
p_value <- mean(null_rho >= observed_rho)
p_value
```

```
[1] 0
```

The observed skill sits well outside the shuffled-data null distribution ($p < 0.01$) — this isn't a pattern 65 noisy points could produce by chance.

Does the direction hold up across time lags?

Per Ye *et al.* (2015), the sharper test for real (vs. spurious/synchrony-driven) causality is whether cross-map skill peaks at a **negative** lag in both directions — meaning each variable's past predicts the other's future, not the reverse:

```
lags <- -4:4


ccm_by_lag <- map_dfr(lags, function(tp) {
  out <- CCM(dataFrame = paramecium_didinium, E = 3, Tp = tp,
             columns = "paramecium", target = "didinium",
             libSizes = "65 65 1", sample = 100, seed = 1, showPlot = FALSE)
  tibble(Tp = tp,
```

```

`paramecium -> didinium` = out$`didinium:paramecium`[1],
`didinium -> paramecium` = out$`paramecium:didinium`[1])
})

ccm_by_lag |>
  pivot_longer(-Tp, names_to = "direction", values_to = "rho") |>
  ggplot(aes(x = Tp, y = rho, color = direction)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2) +
  geom_vline(xintercept = 0, linetype = "dashed", color = "grey50") +
  scale_color_manual(values = c("#1B7837", "#762A83")) +
  labs(x = "Time lag (Tp)", y = expression(rho), color = NULL,
       title = "Cross-map skill across time lags") +
  theme_bw(base_size = 14) +
  theme(legend.position = "bottom")

```



Both directions peak at negative lags — each species' past state predicts the other's future, not the reverse. That's the signature of genuine bidirectional causation, not one variable being dragged along by strong one-way forcing.

Take-home

- CCM tests causality between dynamically coupled variables by asking whether one variable's reconstructed manifold can “recover” the other — and whether that skill **converges** with more data.
- The direction is the confusing part: evidence for $X \rightarrow Y$ comes from Y 's **manifold predicting** X , not the other way round.
- You don't need to add covariates the way regression does — coupling is already embedded in each variable's own dynamics.
- You *do* still need to watch for synchrony (strong one-way forcing mimicking bidirectionality) and shared external drivers (like season) — and the fix is checking convergence and lag structure, not adding a covariate list (Ye *et al.* 2015).
- Real, reciprocal predator-prey coupling (like *Paramecium/Didinium*) should — and does — show convergent CCM in **both** directions.

Further reading

- Ye *et al.* (2015) — the lag-scanning method used above, applied in the original paper to this same predator-prey system plus ice-core CO₂/temperature and sardine/anchovy/SST data.
- [Time Series Analysis Handbook, ch. 6 — Convergent Cross Mapping](#) — more worked examples, including where CCM correctly finds *no* causality (e.g. spurious correlations).
- [Does convergent cross mapping require you to control for other variables? \(CrossValidated\)](#) — the discussion this notebook's “Does CCM need to control for other variables?” section is answering.
- Chang *et al.* (2017) — the **rEDM** walkthrough referenced in [the first notebook](#).

Chang, C.-W., Ushio, M. & Hsieh, C. (2017). [Empirical dynamic modeling for beginners](#). *Ecol Res*, 32, 785–796.

Ye, H., Deyle, E.R., Gilarranz, L.J. & Sugihara, G. (2015). [Distinguishing time-delayed causal interactions using convergent cross mapping](#). *Sci Rep*, 5, 14750.