

Error measurement simulations

Leonardo Capitani

2026-07-17

```
library(brms)
library(tidybayes)
library(tidyverse)
library(ggpubr)

# Make random things reproducible
set.seed(1234)

# Bayes stuff
# Use the cmdstanr backend for Stan because it's faster and more modern than
# the default rstan. You need to install the cmdstanr package first
# (https://mc-stan.org/cmdstanr/) and then run cmdstanr::install_cmdstan() to
# install cmdstan on your computer.
options(mc.cores = 4, # Use 4 cores
        brms.backend = "cmdstanr")
bayes_seed <- 1234

# Colors from MetBrewer
clrs <- MetBrewer::met.brewer("Java")

# Custom ggplot themes to make pretty plots
# Get Roboto Condensed at https://fonts.google.com/specimen/Roboto+Condensed
# Get Roboto Mono at https://fonts.google.com/specimen/Roboto+Mono
theme_pred <- function() {
  theme_minimal(base_family = "Roboto Condensed") +
    theme(panel.grid.minor = element_blank(),
          plot.background = element_rect(fill = "white", color = NA),
          plot.title = element_text(face = "bold"),
          strip.text = element_text(face = "bold"),
          strip.background = element_rect(fill = "grey80", color = NA),
```

```

    axis.title.x = element_text(hjust = 0),
    axis.title.y = element_text(hjust = 0),
    legend.title = element_text(face = "bold"))
}

theme_pred_dist <- function() {
  theme_pred() +
    theme(plot.title = element_markdown(family = "Roboto Condensed", face = "plain"),
          plot.subtitle = element_text(family = "Roboto Mono", size = rel(0.9), hjust = 0),
          axis.text.y = element_blank(),
          panel.grid.major.y = element_blank(),
          panel.grid.minor.y = element_blank())
}

theme_pred_range <- function() {
  theme_pred() +
    theme(plot.title = element_markdown(family = "Roboto Condensed", face = "plain"),
          plot.subtitle = element_text(family = "Roboto Mono", size = rel(0.9), hjust = 0),
          panel.grid.minor.y = element_blank())
}

update_geom_defaults("text", list(family = "Roboto Condensed", lineheight = 1))

```

Introduction

From Gelman et al. 2020 *Regression and other stories*, p. 458 and Gelman's blog [post](#) on 14.04.2024:

Another form of incomplete data is measurement error. In the linear regression $y = a + bx + \text{error}$, measurement error in y does not pose a problem, as it folds right into the error term. Suppose we would like to estimate the “underlying” model $y_i = a + bx_i + e_i$ but y is itself measured with error, so that what we actually observe is $y_i^* = y_i + \eta_i$. In that case, we can combine the two models to get $y_i^* = a + bx_i + e_i + \eta_i$, and if the two error terms are independent and each with mean 0, this reduces to a linear regression of y^* on x , which can be fit directly. **So as long as e_i is independent of η_i , you can just combine them into a single error term.**

Simulate

First , we are going to generate fake data, following the linear relationship $y_i = a + bx_i + e_i$ with a priori fixed parameters values:

```
#simulation for measurement error

set.seed(123)

n <- 1000 # number of observations

x <- runif(n, 0, 10) # predictor x with mean = 0 and sd = 10

a <- 0.2 # intercept, the mean value when x = 0
b <- 0.3 # influence of x on y
sigma <- 0.5 # also call in our model "e". It represents to what extent our model is able to

y <- rnorm(n, a + b*x, sigma)

fake <- data.frame(x,y)
```

Then fit the linear model $y_i = a + bx_i + e_i$ to the fake data:

```
get_prior(formula = y ~ 1 + x, data = fake )
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	tag
	(flat)	b								
	(flat)	b	x							
student_t(3, 1.7, 2.5)	Intercept									
student_t(3, 0, 2.5)	sigma							0		
source										
default										
(vectorized)										
default										
default										

```
fit_1 <- brm(formula = y ~ 1 + x,
             family = gaussian,
             data = fake,
             prior = c(prior(normal(0, 0.2), class = Intercept),
                       prior(normal(0, 0.5), class = b),
```

```

      prior(exponential(1), class = sigma)),
      iter = 2000, warmup = 1000, cores = 4, chains = 4,
      backend = 'cmdstanr',
      silent = 2, refresh = 0,
      seed = 15)

print(fit_1)

```

```

Family: gaussian
Links: mu = identity
Formula: y ~ 1 + x
Data: fake (Number of observations: 1000)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
      total post-warmup draws = 4000

```

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.21	0.03	0.15	0.27	1.00	4485	2994
x	0.30	0.01	0.29	0.31	1.00	4779	3011

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	0.50	0.01	0.48	0.52	1.00	3483	3003

Draws were sampled using `sample(hmc)`. For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```

# generate posterior predictions
pred_1 <- fit_1 |>
  add_predicted_draws(newdata = fake, ndraws = 100,
                      seed = 12345) |> mutate(model = "No Measurement error")

```

Note

Take home message:

The model returns exactly the fixed values we have previously determined.

Then we are going to add the error measurement η on y , so $y_i^* = y_i + \eta_i$. We fixed the measurement error as $\eta = 1$. Remember that $e = 0.5$.

```
sigma_y <- 1
fake$y_star <- rnorm(n, fake$y, sigma_y)
```

Then we add some error measurement on x , generated x^* . We fixed the measurement error as $\tau = 4$.

```
sigma_x <- 4
fake$x_star <- rnorm(n, fake$x, sigma_x)
```

Then we fit again the same model, but with error measurement on y

```
get_prior(formula = y_star ~ 1 + x, data = fake )
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	tag
	(flat)	b								
	(flat)	b	x							
student_t(3, 1.7, 2.5)		Intercept								
student_t(3, 0, 2.5)		sigma							0	
source										
default										
(vectorized)										
default										
default										

```
fit_2 <- brm(formula = y_star ~ 1 + x,
  family = gaussian,
  data = fake,
  prior = c(prior(normal(0, 0.2), class = Intercept),
    prior(normal(0, 0.5), class = b),
    prior(exponential(1), class = sigma)),
  iter = 2000, warmup = 1000, cores = 4, chains = 4,
  backend = 'cmdstanr',
  silent = 2, refresh = 0,
  seed = 15)

print(fit_2)
```

```
Family: gaussian
Links: mu = identity
Formula: y_star ~ 1 + x
```

Data: fake (Number of observations: 1000)
 Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
 total post-warmup draws = 4000

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.20	0.07	0.07	0.34	1.00	4166	3134
x	0.29	0.01	0.27	0.32	1.00	4080	3035

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	1.11	0.02	1.07	1.16	1.00	3660	2569

Draws were sampled using `sample(hmc)`. For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
# generate posterior predictions

pred_2 <- fit_2 |>
  add_predicted_draws(newdata = fake, ndraws = 100,
    seed = 12345) |> mutate(model = "Measurement error in y")
```

Then we fit again the same model, but with error measurement on x

```
fit_3 <- brm(formula = y ~ 1 + x_star,
  family = gaussian,
  data = fake,
  prior = c(prior(normal(0, 0.2), class = Intercept),
    prior(normal(0, 0.5), class = b),
    prior(exponential(1), class = sigma)),
  iter = 2000, warmup = 1000, cores = 4, chains = 4,
  backend = 'cmdstanr',
  silent = 2, refresh = 0,
  seed = 15)

print(fit_3)
```

Family: gaussian
 Links: mu = identity
 Formula: y ~ 1 + x_star

Data: fake (Number of observations: 1000)
 Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
 total post-warmup draws = 4000

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	1.15	0.04	1.07	1.22	1.00	4362	3116
x_star	0.10	0.01	0.09	0.11	1.00	6324	3617

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	0.84	0.02	0.81	0.88	1.00	4548	3345

Draws were sampled using `sample(hmc)`. For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
# generate posterior predictions

pred_3 <- fit_3 |>
  add_predicted_draws(newdata = fake, ndraws = 100,
    seed = 12345) |> mutate(model = "Measurement error in x")
```

Then we fit again the same model, but with error measurement on x^* and y^* .

```
fit_4 <- brm(formula = y_star ~ 1 + x_star,
  family = gaussian,
  data = fake,
  prior = c(prior(normal(0, 0.2), class = Intercept),
    prior(normal(0, 0.5), class = b),
    prior(exponential(1), class = sigma)),
  iter = 2000, warmup = 1000, cores = 4, chains = 4,
  backend = 'cmdstanr',
  silent = 2, refresh = 0,
  seed = 15)

print(fit_4)
```

Family: gaussian
 Links: mu = identity
 Formula: y_star ~ 1 + x_star
 Data: fake (Number of observations: 1000)

Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
total post-warmup draws = 4000

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	1.11	0.06	0.99	1.22	1.00	3905	2987
x_star	0.11	0.01	0.09	0.12	1.00	4131	2966

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	1.29	0.03	1.24	1.35	1.00	4181	2789

Draws were sampled using `sample(hmc)`. For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

```
# generate posterior predictions

pred_4 <- fit_4 |>
  add_predicted_draws(newdata = fake, ndraws = 100,
    seed = 12345) |> mutate(model = "Measurement error in both x and y")
```

Merge the four data frames with model's predictions

```
merg <- rbind(pred_1,pred_2,pred_3,pred_4)
```

Then we plot:

```
p1 <- ggplot(pred_1, aes(x = x)) +

  stat_lineribbon(aes(y = .prediction), .width = 0.95, color = "red", fill = "#D1D3D5") +
  geom_point(data = fake, aes(y = y), shape = 1, size = 1, alpha = 0.7) +
  scale_x_continuous(limits = c(-10,20), breaks = seq(from = -10, to = 20, by = 2.5))+
  scale_y_continuous(limits = c(-3, 6))+
  labs(x = "x", y = "y", color = NULL, fill = NULL,
    title = "No measurement error") +
  theme_bw(base_size = 16)

p2 <- ggplot(pred_2, aes(x = x)) +
```



```

stat_lineribbon(aes(y = .prediction), .width = 0.95, color = "red", fill = "#D1D3D5") +
geom_point(data = fake, aes(y = y_star), shape = 1, size = 1, alpha = 0.7) +
labs(x = "x", y = expression(y~"*"),color = NULL, fill = NULL,
      title = "Measurement error on y") +
scale_x_continuous(limits = c(-10,20), breaks = seq(from = -10, to = 20, by = 2.5))+
scale_y_continuous(limits = c(-3, 6))+
theme_bw(base_size = 16)

p3 <- ggplot(pred_3, aes(x = x_star)) +

stat_lineribbon(aes(y = .prediction), .width = 0.95, color = "red", fill = "#D1D3D5") +
geom_point(data = fake, aes(y = y), shape = 1,size = 1, alpha = 0.7) +
labs(x = expression(x~"*"), y = "y",color = NULL, fill = NULL,
      title = "Measurement error on x") +
scale_x_continuous(limits = c(-10,20), breaks = seq(from = -10, to = 20, by = 2.5))+
scale_y_continuous(limits = c(-3, 6))+
theme_bw(base_size = 16)


p4 <- ggplot(pred_4, aes(x = x_star)) +

stat_lineribbon(aes(y = .prediction), , color = "red", fill= "#D1D3D5", .width = 0.95) +
geom_point(data = fake, aes(y = y_star),shape = 1, size = 1, alpha = 0.7) +
labs(x = expression(x~"*"), y =expression(y~"*"),color = NULL, fill = NULL,
      title = "Measurement error in both x and y") +
scale_x_continuous(limits = c(-10,20), breaks = seq(from = -10, to = 20, by = 2.5))+
scale_y_continuous(limits = c(-3, 6))+
theme_bw(base_size = 16)

ggarrange(p1, p2, p3, p4,

          ncol = 2, nrow = 2)

```



We can compare the regression coefficients:

```
tidy_inter_1 <- bayesplot::mcmc_intervals_data(fit_1,
  prob = 0.66, # 66% intervals
  prob_outer = 0.95, # 95% intervals
  pars = vars(-contains(c("prior_", "lprior", "lp_")))) |> mutate(model = "No Measur")

tidy_inter_2 <- bayesplot::mcmc_intervals_data(fit_2,
  prob = 0.66, # 66% intervals
  prob_outer = 0.95, # 95% intervals
  pars = vars(-contains(c("prior_", "lprior", "lp_")))) |> mutate(model = "Measurement error")

tidy_inter_3 <- bayesplot::mcmc_intervals_data(fit_3,
  prob = 0.66, # 66% intervals
  prob_outer = 0.95, # 95% intervals
  pars = vars(-contains(c("prior_", "lprior", "lp_")))) |> mutate(model = "Measurement error")


tidy_inter_4 <- bayesplot::mcmc_intervals_data(fit_4,
  prob = 0.66, # 66% intervals
  prob_outer = 0.95, # 95% intervals
  pars = vars(-contains(c("prior_", "lprior", "lp_")))) |> mutate(model = "Measurement error")
```

```

me <- rbind(tidy_inter_1, tidy_inter_2 , tidy_inter_3, tidy_inter_4) |> filter(parameter != "b_x_star")
levels(me$parameter)[levels(me$parameter) == "b_x_star"] <- "b_x" # rename the b parameter with x
pos <- position_nudge(y = case_when(me$model == "No Measurement error" ~ 0,
                                   me$model == "Measurement error in y" ~ 0.05,
                                   me$model == "Measurement error in x" ~ 0.1,
                                   me$model == "Measurement error in both x and y" ~ 0.15))

ggplot(me, aes(x = m, y = factor(parameter), color = model)) +
  geom_pointinterval(aes(xmin = l, xmax = h), position = pos, size = 8) +
  geom_pointinterval(aes(xmin = ll, xmax = hh), position = pos, size = 2) +
  geom_point(position = pos, color = "black") +
  theme_bw(base_size = 18) +
  scale_color_brewer(palette = "Set1") +
  labs(x = NULL, y = NULL) +
  scale_x_continuous(breaks = seq(from = 0, to = 2, by = 0.15))

```



! Take home messages:

The slope b changes with measurement error in x but not in y .

When you have measurement error in x , two things happen to attenuate b —that is, to pull the regression coefficient toward zero.

1. First, if you spreading out x but keep y unchanged, this will reduce the slope b of y on x .
2. Second, when you add noise to x you're changing the ordering of the data, which will reduce the strength of the relationship.

Fitting a Bayesian regression with error measurements on predictors x

Please see [here](#).